

17s1: COMP9417 Machine Learning and Data Mining

# Unsupervised Learning

May 2, 2017

**Acknowledgement:** Material derived from slides for the book

Machine Learning, Tom M. Mitchell, McGraw-Hill, 1997

<http://www-2.cs.cmu.edu/~tom/mlbook.html>

and slides by Andrew W. Moore available at

<http://www.cs.cmu.edu/~awm/tutorials>

and the book Data Mining, Ian H. Witten and Eibe Frank,

Morgan Kauffman, 2000. <http://www.cs.waikato.ac.nz/ml/weka>

and the book Pattern Classification, Richard O. Duda, Peter E. Hart,

and David G. Stork. Copyright (c) 2001 by John Wiley & Sons, Inc.

and the book “Elements of Statistical Learning”, Trevor Hastie,

Robert Tibshirani and Jerome Friedman. (c) 2001, Springer.

# Aims

This lecture will introduce you to statistical and graphical methods for clustering of “unlabelled” instances in machine learning. Following it you should be able to:

- describe the problem of unsupervised learning
- describe  $k$ -means clustering
- describe the role of the EM algorithm in  $k$ -means clustering
- describe hierarchical clustering
- describe conceptual clustering

Relevant WEKA programs:

`weka.clusterers.EM`, `SimpleKMeans`, `Cobweb`

# Unsupervised vs. Supervised Learning

Informally *clustering* is assignment of objects to classes on basis of observations about objects only, i.e. not given “labels” of the categories of objects by a “teacher”.

**Unsupervised learning** classes initially *unknown* and need to be “discovered” from the data: cluster analysis, class discovery, unsupervised pattern recognition.

**Supervised learning** classes *predefined* and need a “definition” in terms of the data which is used for prediction: classification, discriminant analysis, class prediction, supervised pattern recognition.

## Why unsupervised learning ?

- if labelling expensive, train with small labelled sample then improve with large unlabelled sample
- if labelling expensive, train with large unlabelled sample then learn classes with small labelled sample
- tracking “concept drift” over time by unsupervised learning
- learn new “features” by clustering for later use in classification
- exploratory data analysis with visualization


Note: sometimes the term “classification” is used to mean unsupervised discovery of classes or clusters

# Clustering


- Finding groups of items that are similar
- Clustering is unsupervised
  - The class of an example is not known
- Success of clustering often measured subjectively
  - this is problematic ...
  - there are statistical & other approaches ...
- A data set for clustering is just like a data set for classification, without the class

# Representing clusters

Simple 2-D representation



Venn diagram (Overlapping clusters)




## Representing clusters

Probabilistic assignment

	1	2	3
a	0.4	0.1	0.5
b	0.1	0.8	0.1
c	0.3	0.3	0.4
d	0.1	0.1	0.8
e	0.4	0.2	0.4
f	0.1	0.4	0.5
g	0.7	0.2	0.1
h	0.5	0.4	0.1
...			

Dendrogram





# Cluster analysis

Clustering algorithms form two broad categories: **hierarchical methods** and **partitioning methods**.

Hierarchical algorithms are either **agglomerative** i.e. bottom-up or **divisive** i.e. top-down.

In practice, hierarchical agglomerative methods often used - efficient exact algorithms available.

Partitioning methods usually require specification of no. of clusters, then try to construct the clusters and fit objects to them.

# Representation

Let  $N = \{e_1, \dots, e_n\}$  be a set of elements, i.e. instances.

Let  $\mathcal{C} = (C_1, \dots, C_l)$  be a *partition* of  $N$  into subsets.

Each subset is called a *cluster*, and  $\mathcal{C}$  is called a *clustering*.

Input data can have two forms:

1. each element is associated with a real-valued vector of  $p$  features e.g. measurement levels for different features
2. pairwise similarity data between elements, e.g. correlation, distance (dissimilarity)

Feature-vectors have more information, but similarity is generic (given the appropriate function). Feature-vector matrix:  $N \times p$ , similarity matrix  $N \times N$ . In general, often  $N \gg p$ .

## Clustering framework

The goal of clustering is to find a partition of  $N$  elements (instances) into homogeneous and well-separated clusters. Elements from same cluster should have high similarity, elements from different cluster low similarity. Note: homogeneity and separation not well-defined. In practice, depends on the problem. Also, there are typically interactions between homogeneity and separation - usually, high homogeneity is linked with low separation, and vice versa.

## $k$ -means clustering

Set value for  $k$ , the number of clusters (by prior knowledge or via search)

Initialise: choose points for centres (means) of  $k$  clusters (at random)


Procedure:

1. assign each instance  $x$  to the closest of the  $k$  points
2. re-assign the  $k$  points to be the means of each of the  $k$  clusters
3. repeat 1 and 2 until convergence to a reasonably stable clustering

## Example: one variable 2-means (& standard deviations)

A	51	B	62	B	64	A	48	A	39	A	51
A	43	A	47	A	51	B	64	B	62	A	48
B	62	A	52	A	52	A	51	B	64	B	64
B	64	B	64	B	62	B	63	A	52	A	42
A	45	A	51	A	49	A	43	B	63	A	48
A	42	B	65	A	48	B	65	B	64	A	41
A	46	A	48	B	62	B	66	A	48		
A	45	A	49	A	43	B	65	B	64		
A	45	A	46	A	40	A	46	A	48		

model



## $k$ -means clustering

$P(i)$  is the cluster assigned to element  $i$ ,  $c(j)$  is the centroid of cluster  $j$ ,  $d(v_1, v_2)$  the Euclidean distance between feature vectors  $v_1$  and  $v_2$ .

The goal is to find a partition  $P$  for which the error (distance) function  $E_P = \sum_{i=1}^n d(i, c(P(i)))$  is minimum.

The centroid is the mean or weighted average of the points in the cluster.

$k$ -means very popular clustering tool in many different areas.

Note: can be viewed in terms of the widely-used EM (Expectation-Maximization) algorithm.


## *k*-means clustering algorithm

**Algorithm**      *k*-means

/\* feature-vector matrix  $M(ij)$  is given \*/

1. Start with an arbitrary partition  $P$  of  $N$  into  $k$  clusters
2. for each element  $i$  and cluster  $j \neq P(i)$  let  $E_P^{ij}$  be the cost of a solution in which  $i$  is moved to  $j$ :
  - (a) if  $E_P^{i^*j^*} = \min_{ij} E_P^{ij} < E_P$  then move  $i^*$  to cluster  $j^*$  and repeat step 2 else halt.

## $k$ -means clustering






## $k$ -means clustering

Previous diagram shows three steps to convergence in  $k$ -means with  $k = 3$


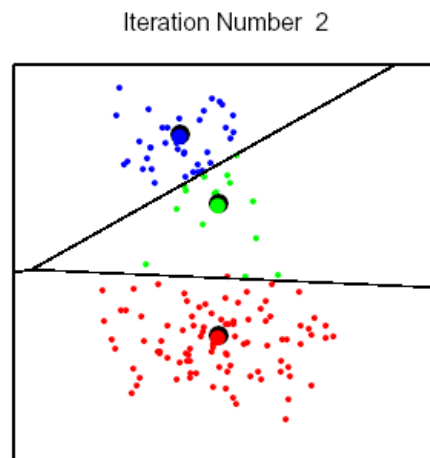


- means move to minimize squared-error criterion
- approximate method of obtaining maximum-likelihood estimates for means
- each point assumed to be in exactly one cluster
- if clusters “blend”, fuzzy  $k$ -means (i.e., overlapping clusters)

Next diagrams show convergence in  $k$ -means with  $k = 3$  for data with two clusters not well separated

## $k$ -means clustering



## $k$ -means clustering




## *k*-means clustering

Trying to *minimize* a loss function in which the goal of clustering is *not* met

- running on microarray data of  $6830 \times 64$  matrix
- total within-cluster sum-of-squares is reduced for  $k = 1$  to 10
- no obvious “correct”  $k$

## $k$ -means clustering



## Practical $k$ -means

- Result can vary significantly based on initial choice of seeds
- Algorithm can get trapped in a local minimum
  - Example: four instances at the vertices of a twodimensional rectangle
    - \* Local minimum: two cluster centers at the midpoints of the rectangle's long sides
- Simple way to increase chance of finding a global optimum: restart with different random seeds
  - can be time-consuming



# Expectation Maximization (EM)

When to use:

- Data is only partially observable
- Unsupervised learning, e.g., clustering (class value “unobservable”)
- Supervised learning (some instance attributes unobservable)

Some uses:

- Train Bayesian Belief Networks
- Unsupervised clustering ( $k$ -means, AUTOCLASS)
- Learning Hidden Markov Models (Baum-Welch algorithm)

## Finite mixtures


Each instance  $x$  generated by

1. Choosing one of the  $k$  Gaussians with uniform probability
2. Generating an instance at random according to that Gaussian

Called *finite mixtures* because there is only a finite number of *generating distributions* being represented.



# Generating Data from Mixture of $k$ Gaussians



## EM for Estimating $k$ Means

Given:

- Instances from  $X$  generated by mixture of  $k$  Gaussian distributions
- Unknown means  $\langle \mu_1, \dots, \mu_k \rangle$  of the  $k$  Gaussians
- Don't know which instance  $x_i$  was generated by which Gaussian

Determine:

- Maximum likelihood estimates of  $\langle \mu_1, \dots, \mu_k \rangle$

## EM for Estimating $k$ Means

Think of full description of each instance as  $y_i = \langle x_i, z_{i1}, z_{i2} \rangle$ , where

- $z_{ij}$  is 1 if  $x_i$  generated by  $j$ th Gaussian, otherwise zero
- $x_i$  observable, from instance set  $x_1, x_2, \dots, x_m$
- $z_{ij}$  unobservable

## EM for Estimating $k$ Means

Initialise: Pick random initial  $h = \langle \mu_1, \mu_2 \rangle$

Iterate:

E step: Calculate expected value  $E[z_{ij}]$  of each hidden variable  $z_{ij}$ , assuming current hypothesis  $h = \langle \mu_1, \mu_2 \rangle$  holds:

$$\begin{aligned} E[z_{ij}] &= \frac{p(x = x_i | \mu = \mu_j)}{\sum_{n=1}^2 p(x = x_i | \mu = \mu_n)} \\ &= \frac{e^{-\frac{1}{2\sigma^2}(x_i - \mu_j)^2}}{\sum_{n=1}^2 e^{-\frac{1}{2\sigma^2}(x_i - \mu_n)^2}} \end{aligned}$$

## EM for Estimating $k$ Means

M step: Calculate new maximum likelihood hypothesis  $h' = \langle \mu'_1, \mu'_2 \rangle$ , assuming value taken on by each hidden variable  $z_{ij}$  is expected value  $E[z_{ij}]$  calculated before. Replace  $h = \langle \mu_1, \mu_2 \rangle$  by  $h' = \langle \mu'_1, \mu'_2 \rangle$ .

$$\mu_j \leftarrow \frac{\sum_{i=1}^m E[z_{ij}] x_i}{\sum_{i=1}^m E[z_{ij}]}$$

i.e.

$$\mu_j \leftarrow \frac{1}{m} \sum_{i=1}^m E[z_{ij}] x_i$$

## EM for Estimating $k$ Means

E step: Calculate probabilities for unknown parameters for each instance

M step: Estimate parameters based on the probabilities

In  $k$ -means the probabilities are stored as instance weights.



## EM Algorithm

Converges to local maximum likelihood  $h$

and provides estimates of hidden variables  $z_{ij}$

In fact, local maximum in  $E[\ln P(Y|h)]$

- $Y$  is complete (observable plus unobservable variables) data
- Expected value taken over possible values of unobserved variables in  $Y$

# General EM Problem

Given:

- Observed data  $X = \{x_1, \dots, x_m\}$
- Unobserved data  $Z = \{z_1, \dots, z_m\}$
- Parameterized probability distribution  $P(Y|h)$ , where
  - $Y = \{y_1, \dots, y_m\}$  is the full data  $y_i = x_i \cup z_i$
  - $h$  are the parameters

Determine:

- $h$  that (locally) maximizes  $E[\ln P(Y|h)]$



## EM for Estimating $k$ Means

Many uses:

- Train Bayesian belief networks
- Unsupervised clustering (e.g.,  $k$  means)
- Hidden Markov Models

## Extending the mixture model

- Using more than two distributions
- Several attributes: easy if independence assumed
- Correlated attributes: difficult
  - Modeled jointly using a bivariate normal distribution with a (symmetric) covariance matrix
  - With  $n$  attributes this requires estimating  $n + n(n + 1)/2$  parameters

## Extending the mixture model

- Nominal attributes: easy if independence assumed
- Correlated nominal attributes: difficult
  - Two correlated attributes result in  $v_1 \times v_2$  parameters
- Missing values: easy
- Distributions other than the normal distribution can be used:
  - “log-normal” if predetermined minimum is given
  - “log-odds” if bounded from above and below
  - Poisson for attributes that are integer counts
- Cross-validation can be used to estimate  $k$  - time consuming !

## General EM Method

Define likelihood function  $Q(h'|h)$  which calculates  $Y = X \cup Z$  using observed  $X$  and current parameters  $h$  to estimate  $Z$

$$Q(h'|h) \leftarrow E[\ln P(Y|h')|h, X]$$

## General EM Method

EM Algorithm:

*Estimation (E) step:* Calculate  $Q(h'|h)$  using the current hypothesis  $h$  and the observed data  $X$  to estimate the probability distribution over  $Y$ .

$$Q(h'|h) \leftarrow E[\ln P(Y|h')|h, X]$$

*Maximization (M) step:* Replace hypothesis  $h$  by the hypothesis  $h'$  that maximizes this  $Q$  function.

$$h \leftarrow \operatorname{argmax}_{h'} Q(h'|h)$$



## Hierarchical clustering

- Bottom up: at each step join the two closest clusters (starting with single-instance clusters)
  - Design decision: distance between clusters
    - \* E.g. two closest instances in clusters vs. distance between means
- Top down: find two clusters and then proceed recursively for the two subsets
  - Can be very fast
- Both methods produce a dendrogram (tree of “clusters”)

## Hierarchical clustering

**Algorithm** Hierarchical agglomerative

/\* dissimilarity matrix  $D(ij)$  is given \*/

1. Find minimal entry  $d_{ij}$  in  $D$  and merge clusters  $i$  and  $j$
2. Update  $D$  by deleting column  $i$  and row  $j$ , and adding new row and column  $i \cup j$
3. Revise entries using  $d_{k,i \cup j} = d_{i \cup j,k} = \alpha_i d_{ki} + \alpha_j d_{kj} + \gamma |d_{ki} - d_{kj}|$
4. If there is more than one cluster then go to step 1.

## Hierarchical clustering

The algorithm relies on a general updating formula. With different operations and coefficients, many different versions of the algorithm can be used to give variant clusterings.

**Single linkage**  $d_{k,i \cup j} = \min(d_{ki}, d_{kj})$  and  $\alpha_i = \alpha_j = \frac{1}{2}$  and  $\gamma = -\frac{1}{2}$ .

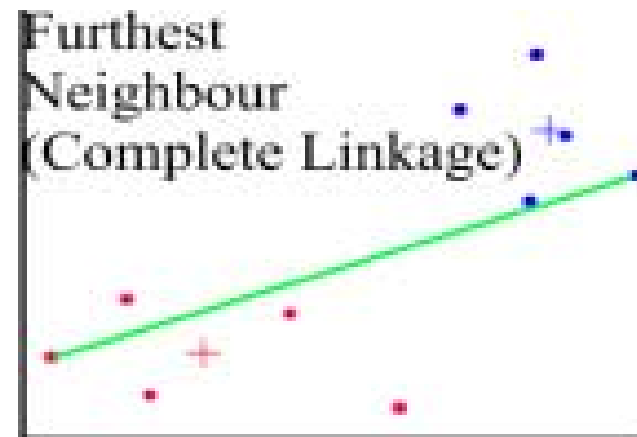



**Complete linkage**  $d_{k,i \cup j} = \max(d_{ki}, d_{kj})$  and  $\alpha_i = \alpha_j = \frac{1}{2}$  and  $\gamma = \frac{1}{2}$ .

**Average linkage**  $d_{k,i \cup j} = \frac{n_i d_{ki}}{n_i + n_j} + \frac{n_j d_{kj}}{n_i + n_j}$  and  $\alpha_i = \frac{n_i}{n_i + n_j}$ ,  $\alpha_j = \frac{n_j}{n_i + n_j}$  and  $\gamma = 0$ .

Note: dissimilarity computed for every pair of points with one point in the first cluster and the other in the second.



## Hierarchical clustering




## Hierarchical clustering

Represent results of hierarchical clustering with a *dendrogram*


See next diagram

- at level 1 all points in individual clusters
- $x_6$  and  $x_7$  are most similar and are merged at level 2
- dendrogram drawn to scale to show similarity between grouped clusters

## Hierarchical clustering



## Hierarchical clustering



Alternative representation of hierarchical clustering based on sets shows hierarchy but not distance

# Dendrograms

Two things to beware of:

1. tree structure is not unique for given clustering - for each bottom-up merge the sub-tree to the right or left must be specified -  $2^{n-1}$  ways to permute the  $n$  leaves in a dendrogram
2. hierarchical clustering imposes a bias - the clustering forms a dendrogram despite the possible lack of an implicit hierarchical structuring in the data


## Dendrograms

Next diagram: average-linkage hierarchical clustering of microarray data


Followed by:

- average-linkage based on average dissimilarity between groups
- complete-linkage based on dissimilarity of furthest pair between groups
- single-linkage based on dissimilarity of closest pair between groups

# Dendrograms




## Dendrograms





## Dendrograms





## Conceptual clustering

- COBWEB/CLASSIT: incrementally forms a hierarchy of clusters (nominal/numerical attributes)
- In the beginning tree consists of empty root node
- Instances are added one by one, and the tree is updated appropriately at each stage
- Updating involves finding the right leaf for an instance (possibly restructuring the tree)
- Updating decisions are based on *category utility*

## Category utility

Category utility is a kind of quadratic loss function defined on conditional probabilities:

$$CU(C_1, C_2, \dots, C_k) = \frac{\sum_l \Pr[C_l] (\sum_i \sum_j \Pr[a_i = v_{ij} \mid C_l]^2 - \Pr[a_i = v_{ij}]^2)}{k}$$

where

- $C_1, C_2, \dots, C_k$  are the  $k$  clusters
- $a_i$  is the  $i$ th attribute with values  $v_{i1}, v_{i2}, \dots$
- intuition: knowing class  $C_l$  gives a better estimate of values of attributes than not knowing it
- measure amount by which that knowledge helps in the probability estimates

## Category utility

Division by  $k$  prevents overfitting, because

- If every instance gets put into a different category  $\Pr[a_i = v_{ij} \mid C_l] = 1$  for attribute-value in the instance and 0 otherwise
- the numerator becomes ( $m$  = total no. of values for set of attributes):

$$m - \sum_i \sum_j \Pr[a_i = v_{ij}]^2$$

- and division by  $k$  penalizes large numbers of clusters

## Category utility

Category utility can be extended to numerical attributes by assuming normal distribution on attribute values.

- estimate standard deviation of attributes and use in formula
- impose minimum variance threshold as a heuristic



## Probability-based clustering

- Problems with above heuristic approach:
  - Division by  $k$ ?
  - Order of examples?
  - Are restructuring operations sufficient?
  - Is result at least local minimum of category utility?
- From a probabilistic perspective, we want to find the most likely clusters given the data
- Also: instance only has certain probability of belonging to a particular cluster

## MDL and clustering

- Description length (DL) needed for encoding the clusters (e.g. cluster centers)
- DL of data given theory: need to encode cluster membership and position relative to cluster (e.g. distance to cluster center)
- Works if coding scheme needs less code space for small numbers than for large ones
- With nominal attributes, we need to communicate probability distributions for each cluster

## Bayesian clustering

- Problem: overfitting possible if number of parameters gets large
- Bayesian approach: every parameter has a prior probability distribution
  - Gets incorporated into the overall likelihood figure and thereby penalizes introduction of parameters
- Example: Laplace estimator for nominal attributes
- Can also have prior on number of clusters!
- Actual implementation: NASA's AUTOCLASS
  - P. Cheeseman - recently with NICTA



## Semi-supervised Learning

Problem: obtaining labelled examples may be difficult, expensive

However, may have many unlabelled instances (e.g., documents)

## Semi-supervised Learning

1. Learn initial classifier using labelled set
2. Apply classifier to unlabelled set
3. Learn new classifier from now-labelled data
4. Repeat until convergence

## Self-training algorithm

Given: labelled data  $\langle x, y \rangle$  and unlabelled data  $\langle x \rangle$

Repeat:

Train classifier  $h$  from labelled data using supervised learning

Label unlabelled data using classifier  $h$

Assumes: classifications by  $h$  will tend to be correct (especially high probability ones)

## Example: use Naive Bayes algorithm

Apply self-training algorithm using Naive Bayes

A form of EM training . . .

# Co-training

Blum & Mitchell (1998)

Key idea: two views of an instance,  $f_1$  and  $f_2$

- assume  $f_1$  and  $f_2$  independent and compatible
- if we have a good attribute set, leverage similarity between attribute values in each view, assuming they predict the class, to classify the unlabelled data

## Co-training

Multi-view learning

Given two (or more) perspectives on data, e.g., different attribute sets

Train separate models for each perspective on small set of labelled data

Use models to label a subset of the unlabelled data

Repeat until no more unlabelled examples

## Clustering summary

- many techniques available – may not be single “magic bullet” rather different techniques useful for different aspects of data
- hierarchical clustering gives a view of the complete structure found, without restricting the no. of clusters, but can be computationally expensive
- different linkage methods can produce very different dendrograms
- higher nodes can be very heterogeneous
- problem may not have a “real” hierarchical structure



## Clustering summary

- $k$ -means and SOM avoid some of these problems, but also have drawbacks
- cannot extract “intermediate features” e.g. a subset of features in which a subset of objects is co-expressed
- for all of these methods, can cluster objects or features, but not both together (coupled two-way clustering)
- should all the points be clustered ? modify algorithms to allow points to be discarded
- visualization is important: dendrograms and SOMs are good but further improvements would help



## Clustering summary

- how can the quality of clustering be estimated ?
  - if clusters known, measure proportion of disagreements to agreements
  - if unknown, measure homogeneity (average similarity between feature vectors in a cluster and the centroid) and separation (weighted average similarity between cluster centroids) with aim of increasing homogeneity and decreasing separation
  - silhouette method, etc.
- clustering is only the first step - mainly exploratory; classification, modelling, hypothesis formation, etc.